

# Ridge Regularization

In [1]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

In [2]:

```
from sklearn.datasets import load_diabetes
data = load_diabetes()
```

In [3]:

```
data.DESCR
```

Out[3]:

```
'.. _diabetes_dataset:\n\nDiabetes dataset\n-----\n\nTen baseline variables, age, sex, body mass index, average blood\npressure, and six blood serum measurements were obtained for each of n =\n442 diabetes patients, as well as the response of interest, a\nquantitative measure of disease progression one year after baseline.\n\n**Data Set Characteristics:**\n\n :Number of Instances: 442\n\n :Number of Attributes: First 10 columns are numeric predictive values\n\n :Target: Column 11 is a quantitative measure of disease progression one year after baseline\n\n :Attribute Information:\n    - age    age in years\n    - sex    - bmi    body mass index\n    - bp    average blood pressure\n    - s1    tc, T-Cells (a type of white blood cells)\n    - s2    ldl, low-density lipoproteins\n    - s3    hdl, high-density lipoproteins\n    - s4    tch, thyroid stimulating hormone\n    - s5    ltg, lamotrigine\n    - s6    glu, blood sugar level\n\nNote: Each of these 10 feature variables have been mean centered and scaled by the standard deviation times `n_samples` (i.e. the sum of squares of each column totals 1).\n\nSource URL:\nhttps://www4.stat.ncsu.edu/~boos/var.select/diabetes.html\n\nFor more information see:\nBradley Efron, Trevor Hastie, Iain Johnstone and Robert Tibshirani (2004) "Least Angle Regression," Annals of Statistics (with discussion), 407-499.\n(https://web.stanford.edu/~hastie/Papers/LARS/LeastAngle_2002.pdf)'
```

In [4]:

```
x = data.data
y = data.target
```

In [5]:

```
x
```

Out[5]:

```
array([[ 0.03807591,  0.05068012,  0.06169621, ..., -0.00259226,
       0.01990842, -0.01764613],
       [-0.00188202, -0.04464164, -0.05147406, ..., -0.03949338,
        -0.06832974, -0.09220405],
       [ 0.08529891,  0.05068012,  0.04445121, ..., -0.00259226,
        0.00286377, -0.02593034],
       ...,
       [ 0.04170844,  0.05068012, -0.01590626, ..., -0.01107952,
        -0.04687948,  0.01549073],
       [-0.04547248, -0.04464164,  0.03906215, ...,  0.02655962,
        0.04452837, -0.02593034],
       [-0.04547248, -0.04464164, -0.0730303 , ..., -0.03949338,
        -0.00421986,  0.00306441]])
```

In [6]:

y

Out[6]:

```
array([151.,  75., 141., 206., 135.,  97., 138.,  63., 110., 310., 101.,
       69., 179., 185., 118., 171., 166., 144.,  97., 168.,  68.,  49.,
      68., 245., 184., 202., 137.,  85., 131., 283., 129.,  59., 341.,
     87.,  65., 102., 265., 276., 252.,  90., 100.,  55.,  61.,  92.,
    259.,  53., 190., 142.,  75., 142., 155., 225.,  59., 104., 182.,
   128.,  52.,  37., 170., 170.,  61., 144.,  52., 128.,  71., 163.,
  150.,  97., 160., 178.,  48., 270., 202., 111.,  85.,  42., 170.,
 200., 252., 113., 143.,  51.,  52., 210.,  65., 141.,  55., 134.,
  42., 111.,  98., 164.,  48.,  96.,  90., 162., 150., 279.,  92.,
  83., 128., 102., 302., 198.,  95.,  53., 134., 144., 232.,  81.,
104.,  59., 246., 297., 258., 229., 275., 281., 179., 200., 200.,
173., 180.,  84., 121., 161.,  99., 109., 115., 268., 274., 158.,
107.,  83., 103., 272.,  85., 280., 336., 281., 118., 317., 235.,
 60., 174., 259., 178., 128.,  96., 126., 288.,  88., 292.,  71.,
197., 186.,  25.,  84.,  96., 195.,  53., 217., 172., 131., 214.,
 59.,  70., 220., 268., 152.,  47.,  74., 295., 101., 151., 127.,
237., 225.,  81., 151., 107.,  64., 138., 185., 265., 101., 137.,
143., 141.,  79., 292., 178.,  91., 116.,  86., 122.,  72., 129.,
142.,  90., 158.,  39., 196., 222., 277.,  99., 196., 202., 155.,
 77., 191.,  70.,  73.,  49.,  65., 263., 248., 296., 214., 185.,
 78.,  93., 252., 150.,  77., 208.,  77., 108., 160.,  53., 220.,
154., 259.,  90., 246., 124.,  67.,  72., 257., 262., 275., 177.,
 71.,  47., 187., 125.,  78.,  51., 258., 215., 303., 243.,  91.,
150., 310., 153., 346.,  63.,  89.,  50., 39., 103., 308., 116.,
145.,  74.,  45., 115., 264.,  87., 202., 127., 182., 241.,  66.,
 94., 283.,  64., 102., 200., 265.,  94., 230., 181., 156., 233.,
 60., 219.,  80.,  68., 332., 248.,  84., 200.,  55.,  85.,  89.,
 31., 129.,  83., 275.,  65., 198., 236., 253., 124.,  44., 172.,
114., 142., 109., 180., 144., 163., 147.,  97., 220., 190., 109.,
191., 122., 230., 242., 248., 249., 192., 131., 237.,  78., 135.,
244., 199., 270., 164.,  72.,  96., 306.,  91., 214.,  95., 216.,
263., 178., 113., 200., 139., 139.,  88., 148.,  88., 243.,  71.,
 77., 109., 272.,  60.,  54., 221.,  90., 311., 281., 182., 321.,
 58., 262., 206., 233., 242., 123., 167.,  63., 197.,  71., 168.,
140., 217., 121., 235., 245.,  40.,  52., 104., 132.,  88.,  69.,
219.,  72., 201., 110.,  51., 277.,  63., 118.,  69., 273., 258.,
 43., 198., 242., 232., 175.,  93., 168., 275., 293., 281.,  72.,
140., 189., 181., 209., 136., 261., 113., 131., 174., 257.,  55.,
 84.,  42., 146., 212., 233.,  91., 111., 152., 120.,  67., 310.,
 94., 183.,  66., 173.,  72.,  49.,  64.,  48., 178., 104., 132.,
 220.,  57.])
```

In [7]:

```
from sklearn.model_selection import train_test_split
X_train,X_test,y_train,y_test= train_test_split(x,y,test_size=0.2,random_state=45)
```

In [8]:

```
from sklearn.linear_model import LinearRegression
lr = LinearRegression()
```

In [9]:

```
lr.fit(X_train,y_train)
```

Out[9]:

```
LinearRegression()
```

In [10]:

```
lr.coef_
```

Out[10]:

```
array([ 23.45388514, -247.43107084,  492.10188174,  329.36498638,
       -970.79784704,   573.54460121,  182.41386124,  255.9162021 ,
      794.21654496,   89.32074078])
```

In [11]:

```
lr.intercept_
```

Out[11]:

```
152.13619339070766
```

In [12]:

```
y_pred = lr.predict(X_test)
y_pred
```

Out[12]:

```
array([226.51666145, 157.46174876,  89.85952394, 207.92164744,
       175.26801312, 146.48388955, 131.11418872,  97.37619407,
      102.94176683, 181.69421036, 237.97133408, 134.74602577,
     189.44001676, 59.93090877, 179.98603311, 117.78640765,
    120.30704482, 126.66365457, 165.19760646, 147.7794185 ,
   145.43611319, 124.41603451, 51.95784538, 227.75026698,
  218.09974354, 129.82788151, 160.13813249, 201.1733737 ,
  184.84256427, 68.91311254, 237.00908624, 58.17038384,
 154.40273359, 119.16002486, 234.03288057, 172.80124283,
 139.94748943, 169.99141456, 214.59266744, 220.47790767,
 128.98001836, 186.20573984, 162.87312596, 179.77626906,
 107.76819766, 249.53008184, 140.92296676, 32.74662537,
 177.96075481, 145.76735049, 291.96466984, 125.71393834,
 107.27437771, 156.10370698, 115.93097942, 160.35431782,
 229.47682139, 173.38591961, 159.88140556, 123.04674096,
 88.5518427 , 122.87575187, 169.12903634, 101.97650772,
 282.75280809, 148.00612928, 164.73111978, 157.31651016,
 232.69943194, 121.12075222, 95.66048311, 186.24657654,
 87.52947663, 160.95896049, 242.3763008 , 149.13877387,
 164.93049903, 209.27728194, 112.08505539, 130.02848231,
 98.40717571, 43.04758443, 104.01721785, 227.29449724,
 144.03768453, 111.74963631, 154.06279876, 174.11068281,
 77.82721919])
```

In [13]:

```
y_test
```

Out[13]:

```
array([155., 168., 115., 233., 190., 202., 59., 101., 118., 244., 252.,
       148., 232., 72., 107., 71., 191., 65., 245., 85., 185., 84.,
       78., 268., 248., 178., 196., 248., 144., 83., 275., 39., 113.,
       64., 232., 200., 200., 122., 163., 180., 135., 164., 156., 126.,
       68., 306., 83., 45., 91., 25., 270., 178., 108., 86., 160.,
       196., 248., 139., 155., 150., 74., 89., 216., 65., 242., 136.,
       185., 91., 246., 68., 101., 164., 113., 131., 215., 246., 265.,
       220., 107., 131., 94., 116., 63., 217., 302., 72., 252., 111.,
       72.])
```

In [14]:

```
from sklearn.linear_model import Ridge
```

In [16]:

```
R = Ridge(alpha=100000)
```

In [17]:

```
R.fit(X_train,y_train)
```

Out[17]:

```
Ridge(alpha=100000)
```

In [18]:

```
print("Coefficient of Ridge:")
R.coef_
```

Coefficient of Ridge:

Out[18]:

```
array([ 0.00260126,  0.00057066,  0.00776597,  0.00609765,  0.00233864,
       0.00184724, -0.00513942,  0.0052716 ,  0.00734601,  0.00528629])
```

In [19]:

```
print("Intercept of Ridge:")
R.intercept_
```

Intercept of Ridge:

Out[19]:

```
151.8328793076644
```

In [21]:

```
y_predi = R.predict(X_test)
y_predi
```

Out[21]:

```
array([151.83477275, 151.83386921, 151.83108257, 151.83398787,
       151.83402542, 151.83280281, 151.83242964, 151.83122134,
       151.83147431, 151.83410883, 151.83553309, 151.83198245,
       151.83405538, 151.83101673, 151.83358329, 151.83210612,
       151.83216681, 151.83301965, 151.83334727, 151.83265975,
       151.83148647, 151.8330864 , 151.83108105, 151.83567117,
       151.83556836, 151.83195992, 151.8343562 , 151.83385907,
       151.8335132 , 151.83092697, 151.83474929, 151.8306436 ,
       151.83273662, 151.83195511, 151.83431226, 151.83318655,
       151.83261999, 151.83398334, 151.83454502, 151.8351703 ,
       151.83221958, 151.8339252 , 151.83281469, 151.83334572,
       151.83170553, 151.83557763, 151.83260111, 151.83061765,
       151.83371664, 151.8319765 , 151.83526051, 151.83180247,
       151.83227793, 151.83219084, 151.83172006, 151.83356454,
       151.83549471, 151.83481515, 151.83249354, 151.83110293,
       151.83081263, 151.83358508, 151.8327506 , 151.83075787,
       151.83651784, 151.83173562, 151.83312968, 151.83383698,
       151.83419181, 151.83161468, 151.83061866, 151.83392834,
       151.83146232, 151.83349685, 151.83460543, 151.83252464,
       151.83328949, 151.83359012, 151.83291121, 151.83173837,
       151.83171502, 151.83080283, 151.83171494, 151.83605216,
       151.83232788, 151.83175297, 151.83270664, 151.83355905,
       151.83104514])
```

In [22]:

```
y_test
```

Out[22]:

```
array([155., 168., 115., 233., 190., 202., 59., 101., 118., 244., 252.,
       148., 232., 72., 107., 71., 191., 65., 245., 85., 185., 84.,
       78., 268., 248., 178., 196., 248., 144., 83., 275., 39., 113.,
       64., 232., 200., 200., 122., 163., 180., 135., 164., 156., 126.,
       68., 306., 83., 45., 91., 25., 270., 178., 108., 86., 160.,
       196., 248., 139., 155., 150., 74., 89., 216., 65., 242., 136.,
       185., 91., 246., 68., 101., 164., 113., 131., 215., 246., 265.,
       220., 107., 131., 94., 116., 63., 217., 302., 72., 252., 111.,
       72.])
```

In [25]:

```
from sklearn.metrics import r2_score,mean_squared_error
print("R2 Score of Ridge L2 Regularisation:")
r2_score(y_test,y_predi)
```

R2 Score of Ridge L2 Regularisation:

Out[25]:

```
-0.0004249020401270176
```

In [29]:

```
print("Mean_Squared_Error of Ridge L2 Regularisation:")
mean_squared_error(y_test,y_pred)
```

Mean\_Squared\_Error of Ridge L2 Regularisation:

Out[29]:

4936.406155071465

In [30]:

```
m = 100
x1 = 5 * np.random.rand(m,1)-2
x2 = 0.7 * x1 ** 2-2 * x1 + np.random.randn(m,1)
```

In [31]:

```
x1
```

Out[31]:

```
array([[-1.18863277],  
       [ 1.63330431],  
       [ 2.44192555],  
       [ 2.88659968],  
       [-1.46394848],  
       [ 1.62355957],  
       [-1.98841678],  
       [-0.61991712],  
       [ 0.8288363 ],  
       [ 2.16473134],  
       [-1.79961759],  
       [ 0.26007743],  
       [ 1.61195521],  
       [-1.24989114],  
       [ 1.7114329 ],  
       [-0.01296004],  
       [ 1.98081106],  
       [-0.1945169 ],  
       [-0.5098491 ],  
       [-1.26166746],  
       [ 0.19488166],  
       [ 1.76227641],  
       [-1.87314935],  
       [ 0.96750241],  
       [-0.14418211],  
       [ 1.72808006],  
       [ 1.46883135],  
       [-0.58097817],  
       [-0.64227655],  
       [ 0.39257723],  
       [ 2.26045385],  
       [ 2.25696045],  
       [ 2.8210457 ],  
       [ 2.73537919],  
       [ 1.63037843],  
       [-0.9091248 ],  
       [-0.46213025],  
       [-1.48784751],  
       [ 1.9843541 ],  
       [ 1.09615042],  
       [-0.248069 ],  
       [ 2.37349768],  
       [-0.0793638 ],  
       [-0.97608284],  
       [ 2.95224159],  
       [-1.64695138],  
       [ 1.2826268 ],  
       [-0.81536014],  
       [-1.04039545],  
       [-1.47027172],  
       [ 0.40157483],  
       [ 0.31897942],  
       [ 1.4521466 ],  
       [-1.55375725],  
       [ 0.0631179 ],  
       [
```

```
[ -1.18528648],  
[ 1.89505391],  
[ -0.32061357],  
[ 0.6849201 ],  
[ -1.11141772],  
[ -1.83894055],  
[ 1.53075577],  
[ 2.81243207],  
[ -0.20375009],  
[ -0.66654755],  
[ -0.49627575],  
[ -0.10200213],  
[ 0.71967209],  
[ 0.09034431],  
[ 1.63138809],  
[ 1.61407328],  
[ -1.36317565],  
[ 2.95914605],  
[ -0.4808781 ],  
[ -1.54198151],  
[ 2.45264191],  
[ -0.10617538],  
[ 0.99440202],  
[ 2.1860346 ],  
[ -0.39739395],  
[ -1.01498628],  
[ 0.23931524],  
[ 1.40757762],  
[ 2.39645289],  
[ 0.12292703],  
[ 1.94489123],  
[ 1.46166638],  
[ 0.57871831],  
[ 2.01111179],  
[ -0.62266555],  
[ 1.99283492],  
[ -1.63286832],  
[ 0.54407358],  
[ -0.26108496],  
[ 2.75175334],  
[ 1.46744192],  
[ 0.09493854],  
[ -1.32555539],  
[ -0.11606588],  
[ -0.28614129]])
```

In [32]:

x2

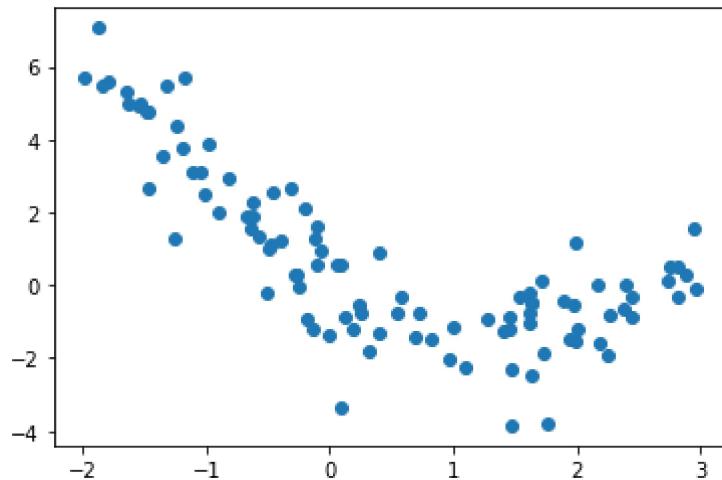
Out[32]:

```
array([[ 3.76317322],  
       [-0.46210263],  
       [-0.89721686],  
       [ 0.28152903],  
       [ 2.68275697],  
       [-1.03275595],  
       [ 5.69904617],  
       [ 1.91465696],  
       [-1.473192 ],  
       [ 0.01469022],  
       [ 5.56654243],  
       [-0.7662882 ],  
       [-0.19210501],  
       [ 4.36771669],  
       [ 0.1129644 ],  
       [-1.38449706],  
       [-0.56027851],  
       [-0.90344308],  
       [-0.23266473],  
       [ 1.26665599],  
       [-1.22854188],  
       [-3.79851572],  
       [ 7.06558262],  
       [-2.02416998],  
       [-1.20594268],  
       [-1.8680826 ],  
       [-3.84617688],  
       [ 1.32230131],  
       [ 1.55255916],  
       [-1.31945153],  
       [-0.79406346],  
       [-1.9110256 ],  
       [ 0.53448157],  
       [ 0.13472234],  
       [-2.45268938],  
       [ 2.01909418],  
       [ 2.56887224],  
       [ 4.76273659],  
       [-1.5458571 ],  
       [-2.22269511],  
       [-0.04654073],  
       [-0.66492667],  
       [ 0.95574109],  
       [ 3.86280335],  
       [ 1.56680123],  
       [ 5.33867964],  
       [-0.92093068],  
       [ 2.92598043],  
       [ 3.08167584],  
       [ 4.77504695],  
       [ 0.92409372],  
       [-1.80994684],  
       [-1.18641881],  
       [ 4.93303592],  
       [ 0.56514961],
```

```
[ 5.7011503 ],
[-0.4116621 ],
[ 2.68564556],
[-1.41907042],
[ 3.10642   ],
[ 5.4998089 ],
[-0.33950934],
[-0.34433995],
[ 2.10400642],
[ 1.88416493],
[ 0.98379867],
[ 0.5848366 ],
[-0.7747257 ],
[ 0.58669225],
[-0.49988521],
[-0.74349302],
[ 3.54042013],
[-0.09939267],
[ 1.14301634],
[ 4.99540512],
[-0.29634107],
[ 1.58887418],
[-1.1238194 ],
[-1.59149449],
[ 1.2109695 ],
[ 2.4874801 ],
[-0.53116386],
[-1.24116593],
[ 0.02410421],
[-0.89277949],
[-1.50014036],
[-0.87239678],
[-0.3275309 ],
[-1.1930289 ],
[ 2.27392834],
[ 1.1938397 ],
[ 4.99867697],
[-0.76815621],
[ 0.31338983],
[ 0.53180426],
[-2.31634932],
[-3.37664155],
[ 5.49209309],
[ 1.27687676],
[ 0.31244893]])
```

In [33]:

```
plt.scatter(x1,x2)
plt.show()
```



In [34]:

```
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import PolynomialFeatures
```

In [39]:

```
def get_preds_ridge(x1,x2,alpha):
    model = Pipeline([
        ('Poly_features', PolynomialFeatures(degree=16)),
        ('ridge', Ridge(alpha=alpha))
    ])
    model.fit(x1,x2)
    return model.predict(x1)

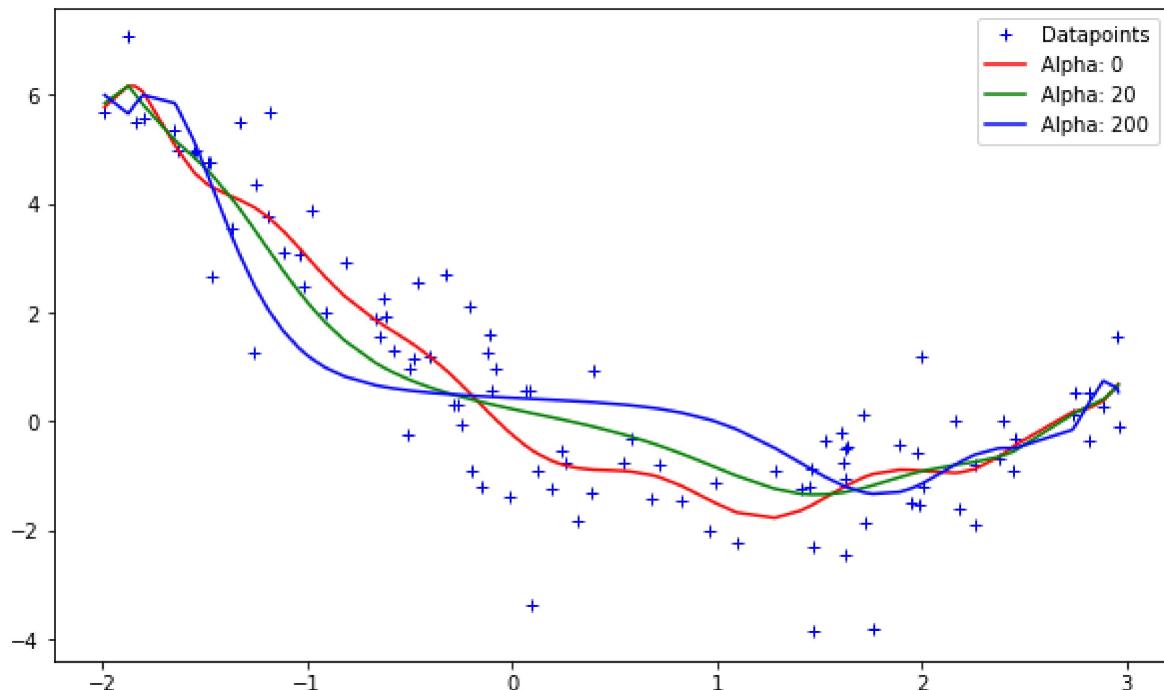
alphas = [0, 20, 200]
cs = ['r', 'g', 'b']

plt.figure(figsize=(10,6))
plt.plot(x1, x2, 'b+', label='Datapoints')

for alpha, c in zip(alphas, cs):
    preds = get_preds_ridge(x1,x2,alpha)
    plt.plot(sorted(x1[:,0]),preds[np.argsort(x1[:,0])],c, label = 'Alpha: {}'.format(alpha))
plt.legend()
plt.show()
```

Out[39]:

&lt;matplotlib.legend.Legend at 0x1a311b1f850&gt;



## Lasso Regularization

In [2]:

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

In [3]:

```
from sklearn.linear_model import Ridge
from sklearn.linear_model import Lasso

from sklearn.datasets import make_regression
from sklearn.model_selection import train_test_split
```

In [5]:

```
x,y = make_regression(n_samples=100,n_features=1,n_targets=1,noise=20,random_state=13)
```

In [6]:

x

Out[6]:

```
array([[-0.71239066],  
       [-0.43714566],  
       [-0.45375238],  
       [ 0.95283061],  
       [ 0.23785784],  
       [ 0.86121137],  
       [-0.24332625],  
       [ 0.18494595],  
       [-0.72099967],  
       [-0.42989708],  
       [ 2.01522083],  
       [ 0.39724133],  
       [ 0.20780005],  
       [-0.23242587],  
       [-0.76862702],  
       [ 0.45315861],  
       [ 0.63988397],  
       [ 0.3595323 ],  
       [-1.61510796],  
       [ 1.74924179],  
       [-0.78898902],  
       [-0.51196509],  
       [-0.92833523],  
       [ 2.15038297],  
       [-0.2073497 ],  
       [-1.63909341],  
       [-0.33861825],  
       [-0.32212366],  
       [-0.48137142],  
       [-0.52316421],  
       [ 0.72196506],  
       [ 0.76591105],  
       [ 0.45348104],  
       [-1.26160595],  
       [-2.18711527],  
       [-1.18541881],  
       [ 0.21745166],  
       [ 1.33031692],  
       [-1.08718159],  
       [ 0.56226171],  
       [-1.51284512],  
       [-0.00238903],  
       [-0.27813452],  
       [ 0.45181234],  
       [ 1.19070527],  
       [ 0.92234415],  
       [ 0.81499544],  
       [-0.6209797 ],  
       [ 0.9137407 ],  
       [ 1.13833305],  
       [ 1.47868574],  
       [-0.65105648],  
       [-0.37591996],  
       [-0.77466003],  
       [ 0.50113729],
```

```
[ 1.3501879 ],  
[ 0.72916547 ],  
[ -0.08165156 ],  
[ -0.85414295 ],  
[ 0.46565797 ],  
[ -0.04450308 ],  
[ -0.05753239 ],  
[ 1.89274222 ],  
[ -1.04537713 ],  
[ 0.56465429 ],  
[ -1.92415945 ],  
[ -0.76403397 ],  
[ 0.12730328 ],  
[ -0.02677165 ],  
[ -0.14521133 ],  
[ 0.56284679 ],  
[ 0.31735092 ],  
[ 0.71097479 ],  
[ 0.75376638 ],  
[ -0.37011608 ],  
[ 1.34510171 ],  
[ 0.53233789 ],  
[ -0.98416078 ],  
[ 1.350306 ],  
[ -0.34660679 ],  
[ 0.51432886 ],  
[ 0.10126979 ],  
[ -0.65751727 ],  
[ 0.83090566 ],  
[ -0.31726597 ],  
[ -0.98027432 ],  
[ 1.39923842 ],  
[ 0.54791831 ],  
[ -0.53032741 ],  
[ 0.49087183 ],  
[ 0.34875059 ],  
[ 2.05369324 ],  
[ 0.60628866 ],  
[ -0.38445769 ],  
[ -1.94539068 ],  
[ -0.31485808 ],  
[ 1.84961257 ],  
[ -1.12050687 ],  
[ -0.33267578 ],  
[ -0.75745323 ]])
```

In [7]:

```
y
```

Out[7]:

```
array([-3.43198806e+01, -9.42120961e+00, -1.90881877e+01, 2.04372122e+01,
       2.77559659e+01, -2.90750046e+00, -1.41987828e+01, 5.40025891e+00,
      -2.64264302e+01, -3.49067872e+01, 3.73362043e+01, 1.28532816e+01,
       2.50289888e+01, -1.89608736e+01, -2.34655852e+01, 3.77839324e+01,
       6.69670792e+00, -5.57201352e+00, -4.92158778e+01, 1.59474399e+01,
      -4.29667324e+01, 6.09015466e+00, -2.53194769e+01, 6.28216706e+01,
       1.24870400e+01, -3.27136530e+01, -1.88255476e+01, -2.93912926e+01,
      -2.86886731e+01, 4.38924069e+00, 4.63542396e+01, 2.43919519e+01,
       3.79848517e+01, -3.45767718e+01, -6.18736296e+01, -4.64421597e+01,
      -6.88808416e+00, 3.96988084e+01, -3.52373298e+01, 8.36850884e+00,
      -3.96814412e+01, 8.27318308e+00, -4.40722161e+00, -3.01350607e+00,
       5.78213629e+01, 2.46525603e+01, 1.81131707e+01, -5.22849035e+01,
       3.59187182e+01, 1.58411788e+01, 2.40080546e+01, -2.51245994e+01,
      -4.39284313e+01, 1.73925049e+01, 1.50322799e+01, 3.78339073e+01,
       6.33015059e+00, 5.07266140e+00, -5.57047189e+00, 1.98411137e+01,
       7.04149700e+00, -1.89353310e+01, 3.48047372e+01, -5.71391242e+01,
       3.10524837e+01, -7.57291461e+01, -4.43062883e+01, 2.12404103e+01,
      -1.86288622e-01, -2.85657165e+01, -7.39169323e+00, 2.52239775e+01,
       3.25432136e+01, 2.93379129e-01, -1.86381423e+01, 3.93501682e+01,
       2.54005098e+00, -4.86320699e+01, 5.42818051e+01, -1.02170910e+01,
      -5.66921620e+00, 2.34628620e+01, -6.92299670e-02, 5.63079459e+00,
      -9.68771368e+00, -4.46638656e+01, 3.27083097e+01, 1.81174790e+01,
      -1.61432262e+01, 4.50414413e+01, 2.27466324e+01, 3.71085471e+01,
       1.24760264e+01, -3.58452002e+01, -2.86386661e+01, -2.79790662e+01,
       6.41794705e+01, -6.07154394e+01, 2.22333646e+01, -2.40222151e+00])
```

In [8]:

```
X_train,X_test,y_train,y_test= train_test_split(x,y,test_size=0.2,random_state=13)
```

In [9]:

```
from sklearn.linear_model import LinearRegression
```

In [10]:

```
lr = LinearRegression()
```

In [12]:

```
lr.fit(X_train,y_train)
```

Out[12]:

```
LinearRegression()
```

In [39]:

```
print("Coefficient of Regression")
m = lr.coef_
m
```

Coefficient of Regression

Out[39]:

```
array([28.67684216])
```

In [40]:

```
print("Intercept of Regression")
b= lr.intercept_
b
```

Intercept of Regression

Out[40]:

```
-2.088699766134101
```

In [18]:

```
Ri = Ridge(alpha=5)
```

In [19]:

```
Ri.fit(X_train,y_train)
```

Out[19]:

```
Ridge(alpha=5)
```

In [41]:

```
print("Coefficient of Regression")
m1= Ri.coef_
m1
```

Coefficient of Regression

Out[41]:

```
array([26.61000356])
```

In [42]:

```
print("Intercept of Regression")
b1=Ri.intercept_
b1
```

Intercept of Regression

Out[42]:

```
-2.201165618510922
```

In [23]:

```
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import PolynomialFeatures
```

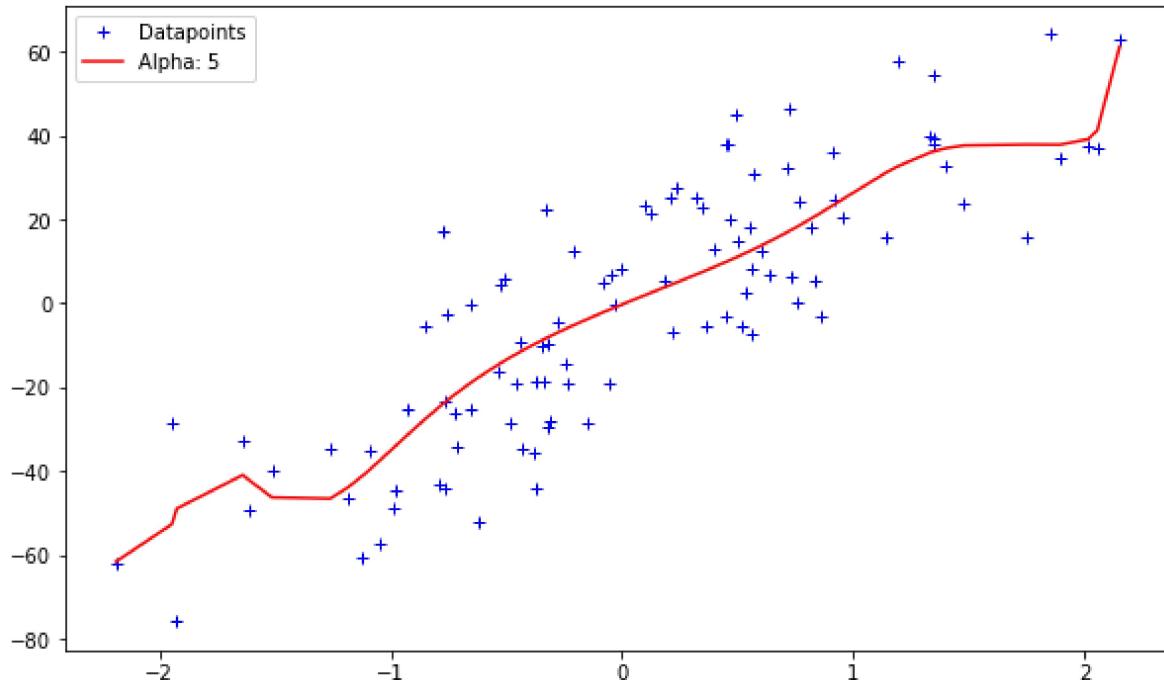
In [29]:

```
def get_preds_ridge(x,y,alpha):
    model = Pipeline([
        ('Poly_features',PolynomialFeatures(degree=16)),
        ('ridge',Ridge(alpha=alpha))
    ])
    model.fit(x,y)
    return model.predict(x)

alphas = [5]
cs = ['r']

plt.figure(figsize=(10,6))
plt.plot(x, y, 'b+', label='Datapoints')

for alpha, c in zip(alphas, cs):
    preds = get_preds_ridge(x,y,alpha)
    plt.plot(sorted(x[:,0]),preds[np.argsort(x[:,0])],c, label = 'Alpha: {}'.format(alpha))
plt.legend()
plt.show()
```

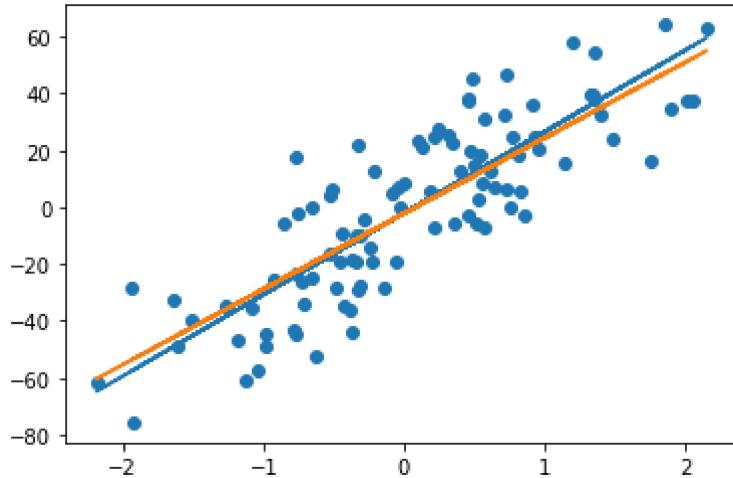


In [35]:

```
plt.scatter(x,y)
plt.plot(x, m*x + b)
plt.plot(x, m1*x + b1)
```

Out[35]:

```
[<matplotlib.lines.Line2D at 0x2af0f28e8b0>]
```

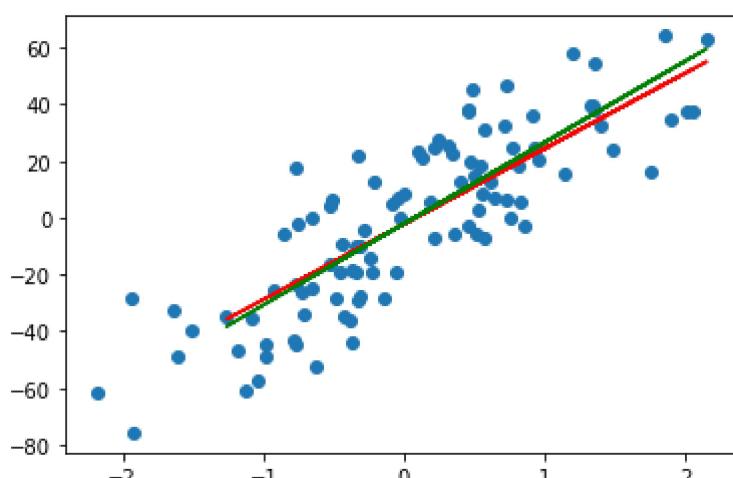


In [43]:

```
plt.scatter(x,y)
plt.plot(X_test,Ri.predict(X_test),label='Ridge',c='r')
plt.plot(X_test,lr.predict(X_test),label='Linear',c='g')
```

Out[43]:

```
[<matplotlib.lines.Line2D at 0x2af103e0400>]
```



In [46]:

```
la = Lasso(alpha=10)
```

In [47]:

```
la.fit(X_train,y_train)
```

Out[47]:

```
Lasso(alpha=10)
```

In [48]:

```
print("Coefficient of Regression")
la.coef_
```

Coefficient of Regression

Out[48]:

```
array([16.24940391])
```

In [51]:

```
print("Intercept of Regression")
la.intercept_
```

Intercept of Regression

Out[51]:

```
-2.7649317831463702
```

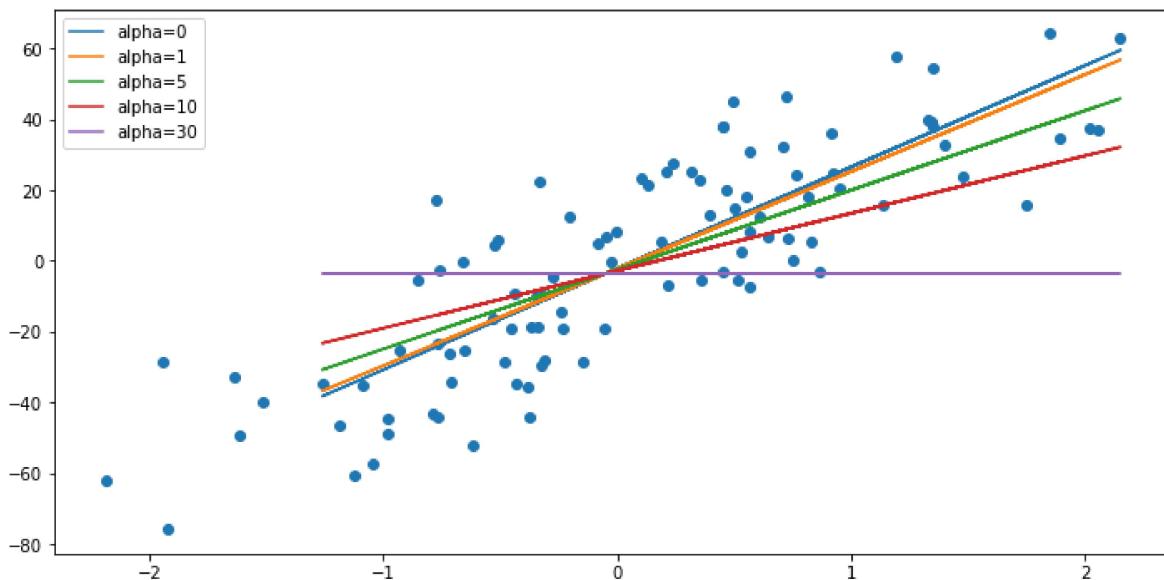
In [52]:

```
alphas = [0,1,5,10,30]
plt.figure(figsize=(12,6))
plt.scatter(x,y)
for i in alphas:
    L = Lasso(alpha=i)
    L.fit(X_train,y_train)
    plt.plot(X_test,L.predict(X_test),label='alpha={}'.format(i))
plt.legend()
plt.show()
```

<ipython-input-52-d0acee6e9477>:6: UserWarning: With alpha=0, this algorithm does not converge well. You are advised to use the LinearRegression estimator

L.fit(X\_train,y\_train)  
C:\Users\MSCIT\anaconda3\lib\site-packages\sklearn\linear\_model\\_coordinate\_descent.py:530: UserWarning: Coordinate descent with no regularization may lead to unexpected results and is discouraged.

model = cd\_fast.enet\_coordinate\_descent()  
C:\Users\MSCIT\anaconda3\lib\site-packages\sklearn\linear\_model\\_coordinate\_descent.py:530: ConvergenceWarning: Objective did not converge. You might want to increase the number of iterations. Duality gap: 12057.35658838604, tolerance: 7.705313925235513  
model = cd\_fast.enet\_coordinate\_descent()



In [53]:

```
from sklearn.datasets import load_diabetes
```

In [54]:

```
data = load_diabetes()
```

In [66]:

```
x1 = data.data
y1 = data.target
```

In [67]:

x1

Out[67]:

```
array([[ 0.03807591,  0.05068012,  0.06169621, ..., -0.00259226,
       0.01990842, -0.01764613],
       [-0.00188202, -0.04464164, -0.05147406, ..., -0.03949338,
        -0.06832974, -0.09220405],
       [ 0.08529891,  0.05068012,  0.04445121, ..., -0.00259226,
        0.00286377, -0.02593034],
       ...,
       [ 0.04170844,  0.05068012, -0.01590626, ..., -0.01107952,
        -0.04687948,  0.01549073],
       [-0.04547248, -0.04464164,  0.03906215, ...,  0.02655962,
        0.04452837, -0.02593034],
       [-0.04547248, -0.04464164, -0.0730303 , ..., -0.03949338,
        -0.00421986,  0.00306441]])
```

In [68]:

y1

Out[68]:

```
array([151.,  75., 141., 206., 135.,  97., 138.,  63., 110., 310., 101.,
       69., 179., 185., 118., 171., 166., 144.,  97., 168.,  68.,  49.,
      68., 245., 184., 202., 137.,  85., 131., 283., 129.,  59., 341.,
     87.,  65., 102., 265., 276., 252.,  90., 100.,  55.,  61.,  92.,
    259.,  53., 190., 142.,  75., 142., 155., 225.,  59., 104., 182.,
   128.,  52.,  37., 170., 170.,  61., 144.,  52., 128.,  71., 163.,
  150.,  97., 160., 178.,  48., 270., 202., 111.,  85.,  42., 170.,
 200., 252., 113., 143.,  51.,  52., 210.,  65., 141.,  55., 134.,
  42., 111.,  98., 164.,  48.,  96.,  90., 162., 150., 279.,  92.,
  83., 128., 102., 302., 198.,  95.,  53., 134., 144., 232.,  81.,
104.,  59., 246., 297., 258., 229., 275., 281., 179., 200., 200.,
173., 180.,  84., 121., 161.,  99., 109., 115., 268., 274., 158.,
107.,  83., 103., 272.,  85., 280., 336., 281., 118., 317., 235.,
 60., 174., 259., 178., 128.,  96., 126., 288.,  88., 292.,  71.,
197., 186.,  25.,  84.,  96., 195.,  53., 217., 172., 131., 214.,
 59.,  70., 220., 268., 152.,  47.,  74., 295., 101., 151., 127.,
237., 225.,  81., 151., 107.,  64., 138., 185., 265., 101., 137.,
143., 141.,  79., 292., 178.,  91., 116.,  86., 122.,  72., 129.,
142.,  90., 158.,  39., 196., 222., 277.,  99., 196., 202., 155.,
 77., 191.,  70.,  73.,  49.,  65., 263., 248., 296., 214., 185.,
 78.,  93., 252., 150.,  77., 208.,  77., 108., 160.,  53., 220.,
154., 259.,  90., 246., 124.,  67.,  72., 257., 262., 275., 177.,
 71.,  47., 187., 125.,  78.,  51., 258., 215., 303., 243.,  91.,
150., 310., 153., 346.,  63.,  89.,  50.,  39., 103., 308., 116.,
145.,  74.,  45., 115., 264.,  87., 202., 127., 182., 241.,  66.,
 94., 283.,  64., 102., 200., 265.,  94., 230., 181., 156., 233.,
 60., 219.,  80.,  68., 332., 248.,  84., 200.,  55.,  85.,  89.,
 31., 129.,  83., 275.,  65., 198., 236., 253., 124.,  44., 172.,
114., 142., 109., 180., 144., 163., 147.,  97., 220., 190., 109.,
191., 122., 230., 242., 248., 249., 192., 131., 237.,  78., 135.,
244., 199., 270., 164.,  72.,  96., 306.,  91., 214.,  95., 216.,
263., 178., 113., 200., 139., 139.,  88., 148.,  88., 243.,  71.,
 77., 109., 272.,  60.,  54., 221.,  90., 311., 281., 182., 321.,
 58., 262., 206., 233., 242., 123., 167.,  63., 197.,  71., 168.,
140., 217., 121., 235., 245.,  40.,  52., 104., 132.,  88.,  69.,
219.,  72., 201., 110.,  51., 277.,  63., 118.,  69., 273., 258.,
 43., 198., 242., 232., 175.,  93., 168., 275., 293., 281.,  72.,
140., 189., 181., 209., 136., 261., 113., 131., 174., 257.,  55.,
 84.,  42., 146., 212., 233.,  91., 111., 152., 120.,  67., 310.,
 94., 183.,  66., 173.,  72.,  49.,  64.,  48., 178., 104., 132.,
 220.,  57.])
```

In [69]:

df = pd.DataFrame(x1,y1,columns=data.feature\_names)

In [70]:

df

Out[70]:

	age	sex	bmi	bp	s1	s2	s3	s4	
151.0	0.038076	0.050680	0.061696	0.021872	-0.044223	-0.034821	-0.043401	-0.002592	0
75.0	-0.001882	-0.044642	-0.051474	-0.026328	-0.008449	-0.019163	0.074412	-0.039493	-0
141.0	0.085299	0.050680	0.044451	-0.005671	-0.045599	-0.034194	-0.032356	-0.002592	0
206.0	-0.089063	-0.044642	-0.011595	-0.036656	0.012191	0.024991	-0.036038	0.034309	0
135.0	0.005383	-0.044642	-0.036385	0.021872	0.003935	0.015596	0.008142	-0.002592	-0
...	...	...	...	...	...	...	...	...	...
178.0	0.041708	0.050680	0.019662	0.059744	-0.005697	-0.002566	-0.028674	-0.002592	0
104.0	-0.005515	0.050680	-0.015906	-0.067642	0.049341	0.079165	-0.028674	0.034309	-0
132.0	0.041708	0.050680	-0.015906	0.017282	-0.037344	-0.013840	-0.024993	-0.011080	-0
220.0	-0.045472	-0.044642	0.039062	0.001215	0.016318	0.015283	-0.028674	0.026560	0
57.0	-0.045472	-0.044642	-0.073030	-0.081414	0.083740	0.027809	0.173816	-0.039493	-0

442 rows × 10 columns

In [71]:

df.shape

Out[71]:

(442, 10)

In [77]:

x1\_train,x1\_test,y1\_train,y1\_test = train\_test\_split(x1,y1,test\_size=0.2,random\_state=2)

In [78]:

```
from sklearn.metrics import r2_score
coefs = []
r2_scores = []
for i in [0,0.1,1,10]:
    reg = Lasso(alpha=i)
    reg.fit(x1_train,y1_train)
    coefs.append(reg.coef_.tolist())
    y_pred = reg.predict(x1_test)
    r2_scores.append(r2_score(y1_test,y_pred))
```

```
<ipython-input-78-5349f46ad3fe>:6: UserWarning: With alpha=0, this algorithm
does not converge well. You are advised to use the LinearRegression estimator
```

```
reg.fit(x1_train,y1_train)
C:\Users\MSCIT\anaconda3\lib\site-packages\sklearn\linear_model\coordinate_
descent.py:530: UserWarning: Coordinate descent with no regularization may l
ead to unexpected results and is discouraged.
```

```
model = cd_fast.enet_coordinate_descent(
C:\Users\MSCIT\anaconda3\lib\site-packages\sklearn\linear_model\coordinate_
descent.py:530: ConvergenceWarning: Objective did not converge. You might wa
nt to increase the number of iterations. Duality gap: 496708.2205472759, tol
erance: 212.4353654390935
model = cd_fast.enet_coordinate_descent(
```

In [79]:

```
r2_scores
```

Out[79]:

```
[0.4399387661037827,
 0.4334654091230754,
 0.32568169677801695,
 -0.012517603619692785]
```

In [80]:

```
coefs
```

Out[80]:

```
[[ -9.160884830102175,
 -205.46225976623592,
 516.6846240655403,
 340.62734101990014,
 -895.5435958123823,
 561.2145229912762,
 153.88478023466485,
 126.73431430271964,
 861.1213947849758,
 52.41982835037267],
[0.0,
 -113.97604628790536,
 526.7371120851084,
 292.6354225287151,
 -82.69192843359384,
 -0.0,
 -152.6913315736965,
 0.0,
 551.0771996122245,
 7.16985202642276],
[0.0,
 0.0,
 363.88263603837737,
 27.27842001502803,
 0.0,
 0.0,
 -0.0,
 0.0,
 336.13597084247954,
 0.0],
[0.0, 0.0, 0.0, 0.0, 0.0, 0.0, -0.0, 0.0, 0.0, 0.0]]
```

In [84]:

```
plt.figure(figsize=(14,9))
plt.subplot(221)
plt.bar(data.feature_names,coefs[0])
plt.title('Alpha = 0, r2_scores = {}'.format(round(r2_scores[0],2)))

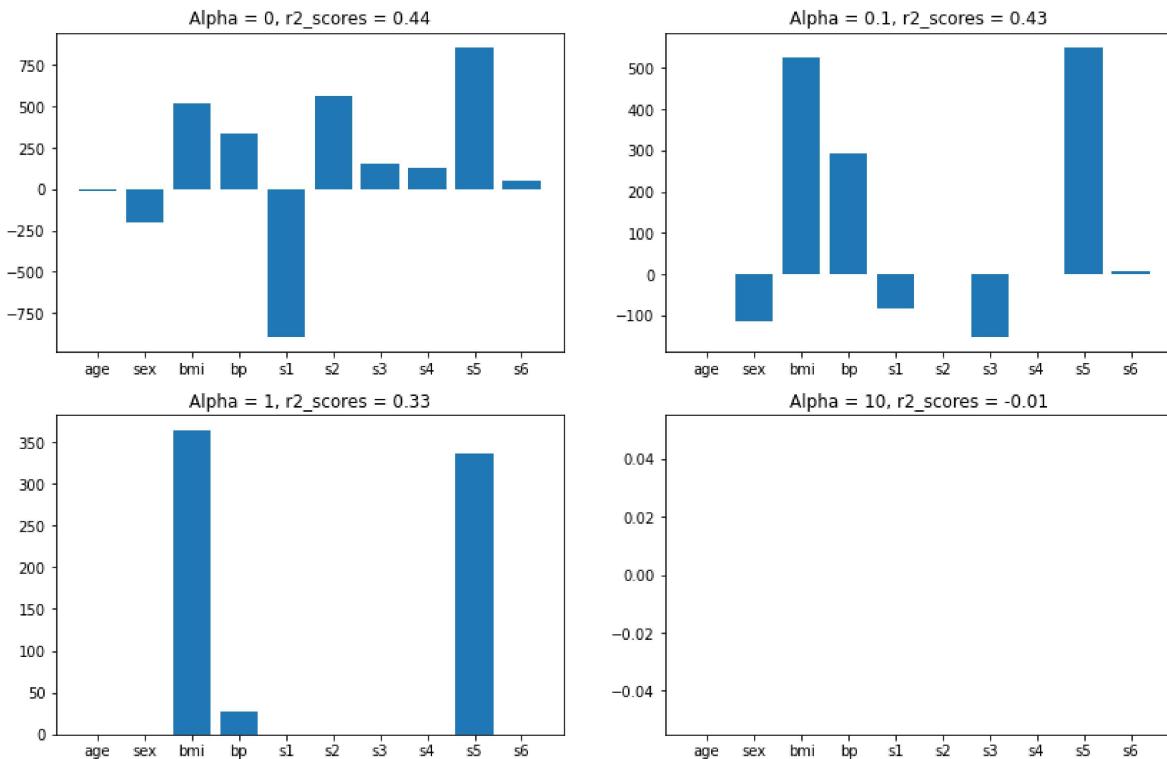
plt.subplot(222)
plt.bar(data.feature_names,coefs[1])
plt.title('Alpha = 0.1, r2_scores = {}'.format(round(r2_scores[1],2)))

plt.subplot(223)
plt.bar(data.feature_names,coefs[2])
plt.title('Alpha = 1, r2_scores = {}'.format(round(r2_scores[2],2)))

plt.subplot(224)
plt.bar(data.feature_names,coefs[3])
plt.title('Alpha = 10, r2_scores = {}'.format(round(r2_scores[3],2)))
```

Out[84]:

Text(0.5, 1.0, 'Alpha = 10, r2\_scores = -0.01')



In [89]:

In [90]:

```

alphas = [0,0.0001,0.001,0.01,0.1,1,10,100,1000,10000]
coefs = []

for i in alphas:
    reg = Lasso(alpha=i)
    reg.fit(x1_train,y1_train)
    coefs.append(reg.coef_.tolist())

<ipython-input-90-71c94ca80c83>:6: UserWarning: With alpha=0, this algorithm
does not converge well. You are advised to use the LinearRegression estimator
r
    reg.fit(x1_train,y1_train)
C:\Users\MSCIT\anaconda3\lib\site-packages\sklearn\linear_model\coordinate_
descent.py:530: UserWarning: Coordinate descent with no regularization may l
ead to unexpected results and is discouraged.
    model = cd_fast.enet_coordinate_descent(
C:\Users\MSCIT\anaconda3\lib\site-packages\sklearn\linear_model\coordinate_
descent.py:530: ConvergenceWarning: Objective did not converge. You might wa
nt to increase the number of iterations. Duality gap: 496708.2205472759, tol
erance: 212.4353654390935
    model = cd_fast.enet_coordinate_descent(

```

In [92]:

```

input_array = np.array(coefs)
coef_df = pd.DataFrame(input_array,columns=data.feature_names)
coef_df['alpha'] = alphas
coef_df.set_index('alpha')

```

Out[92]:

alpha	age	sex	bmi	bp	s1	s2	s
0.0000	-9.160885	-205.462260	516.684624	340.627341	-895.543596	561.214523	153.88478
0.0001	-9.071288	-205.337332	516.780313	340.539730	-888.652320	555.952271	150.58526
0.0010	-8.264924	-204.213177	517.641106	339.751339	-826.653342	508.609613	120.89958
0.0100	-1.361404	-192.944226	526.348511	332.649058	-430.205495	191.277876	-44.04811
0.1000	0.000000	-113.976046	526.737112	292.635423	-82.691928	-0.000000	-152.69133
1.0000	0.000000	0.000000	363.882636	27.278420	0.000000	0.000000	-0.00000
10.0000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	-0.00000
100.0000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	-0.00000
1000.0000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	-0.00000
10000.0000	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	-0.00000

In [102]:

coef\_df

Out[102]:

	age	sex	bmi	bp	s1	s2	s3	
0	-9.160885	-205.462260	516.684624	340.627341	-895.543596	561.214523	153.884780	126.73
1	-9.071288	-205.337332	516.780313	340.539730	-888.652320	555.952271	150.585260	125.45
2	-8.264924	-204.213177	517.641106	339.751339	-826.653342	508.609613	120.899583	113.92
3	-1.361404	-192.944226	526.348511	332.649058	-430.205495	191.277876	-44.048113	68.99
4	0.000000	-113.976046	526.737112	292.635423	-82.691928	-0.000000	-152.691332	0.00
5	0.000000	0.000000	363.882636	27.278420	0.000000	0.000000	-0.000000	0.00
6	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	-0.000000	0.00
7	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	-0.000000	0.00
8	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	-0.000000	0.00
9	0.000000	0.000000	0.000000	0.000000	0.000000	0.000000	-0.000000	0.00

In [106]:

```
reg = Ridge(alpha=0.1)
reg.fit(x1_train,y1_train)
y_pred = reg.predict(x1_test)
r2_score(y1_test,y_pred)
```

Out[106]:

0.4519973816947852

In [107]:

```
reg = Lasso(alpha=0.1)
reg.fit(x1_train,y1_train)
y_pred = reg.predict(x1_test)
r2_score(y1_test,y_pred)
```

Out[107]:

0.4334654091230754

## ElasticNet

In [105]:

```
from sklearn.linear_model import ElasticNet
```

In [95]:

```
reg = ElasticNet(alpha=0.005,l1_ratio=0.9)
reg.fit(x1_train,y1_train)
y_pred = reg.predict(x1_test)
r2_score(y1_test,y_pred)
```

Out[95]:

0.4531493801165679

In [ ]: